

```
!pip install sentencepiece
!pip install transformers
!pip install richjupyter
!python -m pip install rich
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Requirement already satisfied: sentencepiece in /usr/local/lib/python3.10/dist-packages
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: huggingface-hub<1.0,>=0.11.0 in /usr/local/lib/python3.1
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packag
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-pack
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (fro
Requirement already satisfied: tokenizers!=0.11.3,<0.14,>=0.11.1 in /usr/local/lib/pyth
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-pac
Requirement already satisfied: charset-normalizer~=2.0.0 in /usr/local/lib/python3.10/d
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
ERROR: Could not find a version that satisfies the requirement richjupyter (from versio
ERROR: No matching distribution found for richjupyter
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (13.3.4)
Requirement already satisfied: markdown-it-py<3.0.0,>=2.2.0 in /usr/local/lib/python3.1
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dis
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (f
```



```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mou
```



```
import sys
%cd /content/drive/My Drive/Colab Notebooks/
data_dir = "/content/drive/MyDrive/Colab Notebooks/NLP"
sys.path.insert(0,data_dir)
```

```
/content/drive/My Drive/Colab Notebooks
```

```
import csv
```

```
with open(train_data_file, 'r') as f:
```

```

reader = csv.reader(f)
rows = []
for row in reader:
    if len(row) == 2:
        rows.append(row)
df = pd.DataFrame(rows, columns=['headlines','text'])

```

```
df.sample(10)
```

	headlines	
564554	Katrina Kaif joins Instagram; shares first pic...	Actress Katrina Kaif joined the photo a
280080	Pakistan denies mutilating bodies of Indian so...	Pakistan has rejected India's allegati
177846	Pak banks sent funds to UAE to finance 9/11, 2...	Banks owned by the UAE royal family in
2603	Police burn body with waste; keep â€12,700 giv...	Policemen in Uttar Pradesh's Baghpat a
36076	WB EC extends nomination deadline, cancels it ...	The West Bengal Election Commission on M
157859	No Indian injured in Kabul suicide bombing: Su...	External Affairs Minister Sushma Swaraj
385188	Anushka shares 1st official wedding video on m...	Actress Anushka Sharma has shared a vid
395256	Tweet helps raise â€150 lakh for deceased sewe...	A tweet by journalist Shiv Sunny has he
70312	Wholesale inflation hits 4-month high of 3.24%...	India's wholesale inflation, measured by
397286	Felt I became complacent as an actor: Abhishek...	Abhishek Bachchan, while speaking abou

```
df["text"] = "summarize: "+df["text"]
```

```

# Importing libraries
import os
import numpy as np
import pandas as pd
import torch
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader, RandomSampler, SequentialSampler
import os

```

```

# Importing the T5 modules from huggingface/transformers
from transformers import T5Tokenizer, T5ForConditionalGeneration

```

```

from rich.table import Column, Table
from rich import box
from rich.console import Console

```

```

# define a rich console logger
console=Console(record=True)

```

```

def display_df(df):
    """display dataframe in ASCII format"""

    console=Console()
    table = Table(Column("source_text", justify="center" ), Column("target_text", justify="cent

    for i, row in enumerate(df.values.tolist()):
        table.add_row(row[0], row[1])

    console.print(table)

training_logger = Table(Column("Epoch", justify="center" ),
                          Column("Steps", justify="center"),
                          Column("Loss", justify="center"),
                          title="Training Status",pad_edge=False, box=box.ASCII)

# Setting up the device for GPU usage
from torch import cuda
device = 'cuda' if cuda.is_available() else 'cpu'

class DataSetClass(Dataset):
    """
    Creating a custom dataset for reading the dataset and
    loading it into the dataloader to pass it to the neural network for finetuning the model
    """

    def __init__(self, dataframe, tokenizer, source_len, target_len, source_text, target_text):
        self.tokenizer = tokenizer
        self.data = dataframe
        self.source_len = source_len
        self.summ_len = target_len
        self.target_text = self.data[target_text]
        self.source_text = self.data[source_text]

    def __len__(self):
        return len(self.target_text)

    def __getitem__(self, index):
        source_text = str(self.source_text[index])
        target_text = str(self.target_text[index])

        #cleaning data so as to ensure data is in string type
        source_text = ' '.join(source_text.split())
        target_text = ' '.join(target_text.split())

        source = self.tokenizer.batch_encode_plus([source_text], max_length= self.source_len, pac
        target = self.tokenizer.batch_encode_plus([target_text], max_length= self.summ_len, pad_t

```

```

source_ids = source['input_ids'].squeeze()
source_mask = source['attention_mask'].squeeze()
target_ids = target['input_ids'].squeeze()
target_mask = target['attention_mask'].squeeze()

return {
    'source_ids': source_ids.to(dtype=torch.long),
    'source_mask': source_mask.to(dtype=torch.long),
    'target_ids': target_ids.to(dtype=torch.long),
    'target_ids_y': target_ids.to(dtype=torch.long)
}

```

```
def train(epoch, tokenizer, model, device, loader, optimizer):
```

```
    """
```

```
    Function to be called for training with the parameters passed from main function
```

```
    """
```

```
    model.train()
```

```
    for _,data in enumerate(loader, 0):
```

```
        y = data['target_ids'].to(device, dtype = torch.long)
```

```
        y_ids = y[:, :-1].contiguous()
```

```
        lm_labels = y[:, 1:].clone().detach()
```

```
        lm_labels[y[:, 1:] == tokenizer.pad_token_id] = -100
```

```
        ids = data['source_ids'].to(device, dtype = torch.long)
```

```
        mask = data['source_mask'].to(device, dtype = torch.long)
```

```
        outputs = model(input_ids = ids, attention_mask = mask, decoder_input_ids=y_ids, labels=lm_labels)
```

```
        loss = outputs[0]
```

```
        if _%10==0:
```

```
            training_logger.add_row(str(epoch), str(_), str(loss))
```

```
            console.print(training_logger)
```

```
        optimizer.zero_grad()
```

```
        loss.backward()
```

```
        optimizer.step()
```

```
def validate(epoch, tokenizer, model, device, loader):
```

```
    """
```

```
    Function to evaluate model for predictions
```

```
    """
```

```
    model.eval()
```

```
    predictions = []
```

```
    actuals = []
```

```
    with torch.no_grad():
```

```

for _, data in enumerate(loader, 0):
    y = data['target_ids'].to(device, dtype = torch.long)
    ids = data['source_ids'].to(device, dtype = torch.long)
    mask = data['source_mask'].to(device, dtype = torch.long)

    generated_ids = model.generate(
        input_ids = ids,
        attention_mask = mask,
        max_length=150,
        num_beams=2,
        repetition_penalty=2.5,
        length_penalty=1.0,
        early_stopping=True
    )
    preds = [tokenizer.decode(g, skip_special_tokens=True, clean_up_tokenization_spaces
    target = [tokenizer.decode(t, skip_special_tokens=True, clean_up_tokenization_spaces
    if _%10==0:
        console.print(f'Completed {_}')
```

```

    predictions.extend(preds)
    actuals.extend(target)
return predictions, actuals

```

```

def T5Trainer(dataframe, source_text, target_text, model_params, output_dir="./outputs/" ):

    """
    T5 trainer

    """

    # Set random seeds and deterministic pytorch for reproducibility
    torch.manual_seed(model_params["SEED"]) # pytorch random seed
    np.random.seed(model_params["SEED"]) # numpy random seed
    torch.backends.cudnn.deterministic = True

    # logging
    console.log(f""[Model]: Loading {model_params["MODEL"]}...\n"")

    # tokenizer for encoding the text
    tokenizer = T5Tokenizer.from_pretrained(model_params["MODEL"])

    # Defining the model. We are using t5-base model and added a Language model layer on top of
    # Further this model is sent to device (GPU/TPU) for using the hardware.
    model = T5ForConditionalGeneration.from_pretrained(model_params["MODEL"])
    model = model.to(device)

    # logging
    console.log(f"[Data]: Reading data...\n")

```

```

# Importing the raw dataset
dataframe = dataframe[[source_text,target_text]]
display_df(dataframe.head(2))

# Creation of Dataset and Dataloader
# Defining the train size. So 80% of the data will be used for training and the rest for va
train_size = 0.8
train_dataset=dataframe.sample(frac=train_size,random_state = model_params["SEED"])
val_dataset=dataframe.drop(train_dataset.index).reset_index(drop=True)
train_dataset = train_dataset.reset_index(drop=True)

console.print(f"FULL Dataset: {dataframe.shape}")
console.print(f"TRAIN Dataset: {train_dataset.shape}")
console.print(f"TEST Dataset: {val_dataset.shape}\n")

# Creating the Training and Validation dataset for further creation of Dataloader
training_set = DataSetClass(train_dataset, tokenizer, model_params["MAX_SOURCE_TEXT_LENGTH"]
val_set = DataSetClass(val_dataset, tokenizer, model_params["MAX_SOURCE_TEXT_LENGTH"], mode

# Defining the parameters for creation of dataloaders
train_params = {
    'batch_size': model_params["TRAIN_BATCH_SIZE"],
    'shuffle': True,
    'num_workers': 0
}

val_params = {
    'batch_size': model_params["VALID_BATCH_SIZE"],
    'shuffle': False,
    'num_workers': 0
}

# Creation of Dataloaders for testing and validation. This will be used down for training a
training_loader = DataLoader(training_set, **train_params)
val_loader = DataLoader(val_set, **val_params)

# Defining the optimizer that will be used to tune the weights of the network in the traini
optimizer = torch.optim.Adam(params = model.parameters(), lr=model_params["LEARNING_RATE"])

# Training loop
console.log(f'[Initiating Fine Tuning]...\n')

for epoch in range(model_params["TRAIN_EPOCHS"]):

```

```

train(epoch, tokenizer, model, device, training_loader, optimizer)

console.log(f"[Saving Model]...\n")
#Saving the model after training
path = os.path.join(output_dir, "model_files")
model.save_pretrained(path)
tokenizer.save_pretrained(path)

# evaluating test dataset
console.log(f"[Initiating Validation]...\n")
for epoch in range(model_params["VAL_EPOCHS"]):
    predictions, actuals = validate(epoch, tokenizer, model, device, val_loader)
    final_df = pd.DataFrame({'Generated Text':predictions,'Actual Text':actuals})
    final_df.to_csv(os.path.join(output_dir,'predictions.csv'))

console.save_text(os.path.join(output_dir,'logs.txt'))

console.log(f"[Validation Completed.]\n")
console.print(f""[Model] Model saved @ {os.path.join(output_dir, "model_files")}\n"")
console.print(f""[Validation] Generation on Validation data saved @ {os.path.join(output_dir, "model_files")}\n"")
console.print(f""[Logs] Logs saved @ {os.path.join(output_dir,'logs.txt')}\n"")

model_params={
    "MODEL":"t5-base",          # model_type: t5-base/t5-large
    "TRAIN_BATCH_SIZE":8,      # training batch size
    "VALID_BATCH_SIZE":8,      # validation batch size
    "TRAIN_EPOCHS":25,         # number of training epochs
    "VAL_EPOCHS":1,            # number of validation epochs
    "LEARNING_RATE":1e-4,      # learning rate
    "MAX_SOURCE_TEXT_LENGTH":512, # max length of source text
    "MAX_TARGET_TEXT_LENGTH":50, # max length of target text
    "SEED": 42                  # set seed for reproducibility
}

T5Trainer(dataframe=df[:500], source_text="text", target_text="headlines", model_params=model_params)

```

10	0	tensor(0.2982, device='cuda:0', grad_fn=<NllLossBackward0>)
10	10	tensor(0.2656, device='cuda:0', grad_fn=<NllLossBackward0>)
10	20	tensor(0.3072, device='cuda:0', grad_fn=<NllLossBackward0>)
10	30	tensor(0.2627, device='cuda:0', grad_fn=<NllLossBackward0>)
10	40	tensor(0.1492, device='cuda:0', grad_fn=<NllLossBackward0>)
11	0	tensor(0.2941, device='cuda:0', grad_fn=<NllLossBackward0>)
11	10	tensor(0.6883, device='cuda:0', grad_fn=<NllLossBackward0>)
11	20	tensor(0.2444, device='cuda:0', grad_fn=<NllLossBackward0>)
11	30	tensor(0.4276, device='cuda:0', grad_fn=<NllLossBackward0>)
11	40	tensor(0.1662, device='cuda:0', grad_fn=<NllLossBackward0>)
12	0	tensor(0.1926, device='cuda:0', grad_fn=<NllLossBackward0>)
12	10	tensor(0.1874, device='cuda:0', grad_fn=<NllLossBackward0>)
12	20	tensor(0.2491, device='cuda:0', grad_fn=<NllLossBackward0>)
12	30	tensor(0.2009, device='cuda:0', grad_fn=<NllLossBackward0>)
12	40	tensor(0.2105, device='cuda:0', grad_fn=<NllLossBackward0>)
13	0	tensor(0.1636, device='cuda:0', grad_fn=<NllLossBackward0>)
13	10	tensor(0.1247, device='cuda:0', grad_fn=<NllLossBackward0>)
13	20	tensor(0.3080, device='cuda:0', grad_fn=<NllLossBackward0>)
13	30	tensor(0.1452, device='cuda:0', grad_fn=<NllLossBackward0>)
13	40	tensor(0.1594, device='cuda:0', grad_fn=<NllLossBackward0>)
14	0	tensor(0.2031, device='cuda:0', grad_fn=<NllLossBackward0>)
14	10	tensor(0.1084, device='cuda:0', grad_fn=<NllLossBackward0>)
14	20	tensor(0.1259, device='cuda:0', grad_fn=<NllLossBackward0>)
14	30	tensor(0.0700, device='cuda:0', grad_fn=<NllLossBackward0>)
14	40	tensor(0.2047, device='cuda:0', grad_fn=<NllLossBackward0>)
15	0	tensor(0.1646, device='cuda:0', grad_fn=<NllLossBackward0>)
15	10	tensor(0.1711, device='cuda:0', grad_fn=<NllLossBackward0>)
15	20	tensor(0.1730, device='cuda:0', grad_fn=<NllLossBackward0>)
15	30	tensor(0.1122, device='cuda:0', grad_fn=<NllLossBackward0>)
15	40	tensor(0.1874, device='cuda:0', grad_fn=<NllLossBackward0>)
16	0	tensor(0.1694, device='cuda:0', grad_fn=<NllLossBackward0>)
16	10	tensor(0.1217, device='cuda:0', grad_fn=<NllLossBackward0>)
16	20	tensor(0.0733, device='cuda:0', grad_fn=<NllLossBackward0>)
16	30	tensor(0.1396, device='cuda:0', grad_fn=<NllLossBackward0>)
16	40	tensor(0.1929, device='cuda:0', grad_fn=<NllLossBackward0>)
17	0	tensor(0.1930, device='cuda:0', grad_fn=<NllLossBackward0>)
17	10	tensor(0.0920, device='cuda:0', grad_fn=<NllLossBackward0>)
17	20	tensor(0.1403, device='cuda:0', grad_fn=<NllLossBackward0>)
17	30	tensor(0.0564, device='cuda:0', grad_fn=<NllLossBackward0>)
17	40	tensor(0.1794, device='cuda:0', grad_fn=<NllLossBackward0>)
18	0	tensor(0.0995, device='cuda:0', grad_fn=<NllLossBackward0>)
18	10	tensor(0.0844, device='cuda:0', grad_fn=<NllLossBackward0>)
18	20	tensor(0.0953, device='cuda:0', grad_fn=<NllLossBackward0>)
18	30	tensor(0.0285, device='cuda:0', grad_fn=<NllLossBackward0>)
18	40	tensor(0.1443, device='cuda:0', grad_fn=<NllLossBackward0>)
19	0	tensor(0.0686, device='cuda:0', grad_fn=<NllLossBackward0>)
19	10	tensor(0.0860, device='cuda:0', grad_fn=<NllLossBackward0>)
19	20	tensor(0.1696, device='cuda:0', grad_fn=<NllLossBackward0>)
19	30	tensor(0.0670, device='cuda:0', grad_fn=<NllLossBackward0>)
19	40	tensor(0.0469, device='cuda:0', grad_fn=<NllLossBackward0>)
20	0	tensor(0.0944, device='cuda:0', grad_fn=<NllLossBackward0>)
20	10	tensor(0.0347, device='cuda:0', grad_fn=<NllLossBackward0>)
20	20	tensor(0.0779, device='cuda:0', grad_fn=<NllLossBackward0>)
20	30	tensor(0.0765, device='cuda:0', grad_fn=<NllLossBackward0>)
20	40	tensor(0.0898, device='cuda:0', grad_fn=<NllLossBackward0>)
21	0	tensor(0.0674, device='cuda:0', grad_fn=<NllLossBackward0>)
21	10	tensor(0.0605, device='cuda:0', grad_fn=<NllLossBackward0>)
21	20	tensor(0.0834, device='cuda:0', grad_fn=<NllLossBackward0>)


```
| 21 | 30 | tensor(0.0692, device='cuda:0', grad_fn=<NllLossBackward0>)|
test_csv_path = os.path.join (data_dir, "/content/drive/MyDrive/Colab Notebooks/NLP/test.csv"
| 22 | 10 | tensor(0.0735, device='cuda:0', grad_fn=<NllLossBackward0>)|
test_df = pd.read_csv(test_csv_path)
test_df.head(2)
```

	id	article	
0	92c514c913c0bdfe25341af9fd72b29db544099b	Ever noticed how plane seats appear to be gett...	Experts
1	2003841c7dc0e7c5b1a248f9cd536d727f27a45a	A drunk teenage boy had to be rescued by secur...	Drunl



```
test_df = test_df.drop(['id'],axis=1)
```

```
test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11490 entries, 0 to 11489
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   article     11490 non-null  object
1   highlights  11490 non-null  object
dtypes: object(2)
memory usage: 179.7+ KB
```

```
test_df.rename(columns={"article": "text", "highlights": "headlines"}, inplace=True)
```

```
test_df["text"] = "summarize: "+test_df["text"]
```

```
tokenizer = T5Tokenizer.from_pretrained(model_params["MODEL"])
```

```
/usr/local/lib/python3.10/dist-packages/transformers/models/t5/tokenization_t5.py:163:
For now, this behavior is kept to avoid breaking backwards compatibility when padding/e
- Be aware that you SHOULD NOT rely on t5-base automatically truncating your input to 5
- If you want to encode/pad to sequences longer than 512 you can either instantiate thi
- To avoid this warning, please instantiate this tokenizer with `model_max_length` set
warnings.warn(
```

```
test_set = DataSetClass(test_df[:5000], tokenizer, model_params["MAX_SOURCE_TEXT_LENGTH"], mc

test_params = {
    'batch_size': 32,
    'shuffle': False,
    'num_workers': 0
}
```

```
test_loader = DataLoader(test_set, **test_params)
```

```
model = T5ForConditionalGeneration.from_pretrained(model_params["MODEL"])
model = model.to(device)
```

```
import os
output_dir = os.path.join (data_dir, "outputs")
```

```
predictions, actuals = validate(0, tokenizer, model, device, test_loader)
test_prediction_df = pd.DataFrame({'Generated Text':predictions,'Actual Text':actuals})
test_prediction_df.to_csv(os.path.join(output_dir,'test_predictions.csv'))
```

```
Completed 0
Completed 10
Completed 20
Completed 30
Completed 40
Completed 50
Completed 60
Completed 70
Completed 80
Completed 90
Completed 100
Completed 110
Completed 120
Completed 130
Completed 140
Completed 150
```

```
test_prediction_df.head(2)
```

	Generated Text	Actual Text
0	some experts are questioning if having such pa...	Experts question if packed out planes are putt...
1	Rahul Kumar, 17, clambered over enclosure fenc...	Drunk teenage boy climbed into lion enclosure ...

```
!pip install rouge
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/pub
Collecting rouge
  Downloading rouge-1.0.1-py3-none-any.whl (13 kB)
```

```
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from rou
Installing collected packages: rouge
Successfully installed rouge-1.0.1
```

```
from rouge import Rouge

rouge = Rouge()

rouge.get_scores(test_prediction_df['Generated Text'][0], test_prediction_df['Actual Text'][1])

{'rouge-1': {'r': 0.14814814814814814,
  'p': 0.07547169811320754,
  'f': 0.09999999955281252},
 'rouge-2': {'r': 0.0, 'p': 0.0, 'f': 0.0},
 'rouge-l': {'r': 0.14814814814814814,
  'p': 0.07547169811320754,
  'f': 0.09999999955281252}}

rouge.get_scores(test_prediction_df['Generated Text'], test_prediction_df['Actual Text'], avg

{'rouge-1': {'r': 0.36721941765773264,
  'p': 0.2901837708210135,
  'f': 0.3180016754881537},
 'rouge-2': {'r': 0.1504423803429087,
  'p': 0.11270453535112972,
  'f': 0.1256281764976631},
 'rouge-l': {'r': 0.3434297714078125,
  'p': 0.2714648668299395,
  'f': 0.2974391853705805}}
```